

### 3.1

I declared `int invalopcode=0;` in `initialize.c` and put `extern int invalopcode;` in `prototypes.h`.

To modify `_Xint6`, I push registers, compare `invalopcode` with 5, increment, then `iret` until it reaches 5, at which point we call `trap6()`. It prints `invalopcode` and calls `panic()`, hanging the system.

Then, In `main()`, I did `asm("movl %ecx, %cr1");`, an invalid instruction on x86. Because the CPU re-executes that instruction each time, `invalopcode` increments until 5, then `trap6()` is invoked. Because we never return from the invalid instruction or from `panic()`, so no subsequent lines in `main()` run.

#### New Xint6-

```
_Xint6:
    pushal
    cmpl   $5, invalopcode
    jl     .lessThan5
    call   trap6
    jmp     .
.lessThan5:
    incl   invalopcode
    popal
    iret
```

#### Testing-

```
process main(void)
{
    kprintf("About to trigger invalid opcode ... \n");
    asm("movl %ecx, %cr1");
    kprintf("If you see this, #UD was not triggered.\n");
    return OK;
}
```

#### Output-

```
agraw185,Agrawal,Shivli,Spring,2025
About to trigger invalid opcode ...
Trap 6 triggered; invalopcode = 5
```

```
panic: Invalid opcode threshold reached! System halting ...
```

### 3.2

When I ran `asm("int3");` with the default `_Xint3`, I saw XINU print a trap message (exception #3) and then panic. And never reached the line after `int3` in `main()`. So I replaced `_Xint3`'s body with a single `iret` and saw both 'Before `int3`' and 'After `int3`', no trap message, because `_Xint3` immediately returned to the next instruction.

This implies that the CPU saves on the stack the address of the instruction immediately following the `int3`. Because when debuggers use a breakpoint instruction, they want to resume execution after the `int3` once the debugger is done. So by design, the CPU sets EIP in such a way that an `iret` effectively steps over the breakpoint. This shows that the CPU saves the address of the next instruction on the stack for breakpoint exceptions, letting `iret` return to the instruction after the `int3`.

Testing-

```
kprintf("Before int3\n");
asm("int3");
kprintf("After int3\n");
```

Output-

```
agraw185,Agrawal,Shivli,Spring,2025
Before int3
After int3
```

### 3.3

On replacing `asm("movl %ecx, %cr1");` with `asm("int $6")` I saw that `int $6` triggers the same exception #6, but the CPU pushes the address of the next instruction on the stack. If the `_Xint6` only did `iret`, we would return after the `int $6` instruction and continue in `main()` rather than re-executing the invalid op. On using an actually invalid opcode like `movl %ecx, %cr1`, the CPU saves the address of the same failing instruction, so it re-faults on `iret`. So the software interrupt instructions save EIP pointing past the instruction which is not like hardware fault instructions that re-try the same failing instruction.

Then, I used `asm("int3");` to generate exception #3. I replaced it with `asm("int $3");`. Both end up calling `_Xint3`.

`int3` is a 1-byte opcode (0xCC) and `int $3` is a 2-byte sequence (0xCD 0x03).

The CPU's behavior is the same for #3 in that it sets EIP to the next instruction on the stack. If `_Xint3` does an `iret`, it resumes after the `int $3`. Debuggers insert the single-byte `int3` (0xCC) at runtime into user code for breakpoints because it's easier to overwrite just one byte.

Now,, I wrote `main()` so it prints "Before start of processing loop", then enters a `while(1)` loop. Inside the loop, it checks the global `clktime` (in seconds). Once it reaches 8, we print "8 second mark reached" and call `exit()`. Initially, it takes almost 8 real-world seconds to reach `clktime = 8`.

Then, I added `asm("int $32");` at the top of the loop before checking `clktime`. That code often triggered the clock interrupt handler more frequently than the hardware's 1 kHz timer alone. As a result, `clktime` incremented faster in software, so the 8 second mark could be reached in much less than 8 real seconds. We understood that repeatedly forcing interrupt #32 in software can desynchronize `clktime` from real-world time.

## Testing-

```
process main(void)
{
    kprintf("Before start of processing loop\n");
    while (1)
    {
        asm("int $32"); // test with and without

        if (clktime ≥ 8) {
            kprintf("8 second mark reached\n");
            exit();
        }
    }
    return OK;
}
```

## Output-

### first terminal output-

```
agraw185,Agrawal,Shivli,Spring,2025
Using int $3 instead of int3
Back from int $3
```

### Second-

```
agraw185,Agrawal,Shivli,Spring,2025
Before start of processing loop
8 second mark reached (in actually 8 seconds)
```

### Third-

```
agraw185,Agrawal,Shivli,Spring,2025
Before start of processing loop
8 second mark reached (the mark reached very quickly this time though)
```

## 3.4

Generating Interrupt #55, I placed `asm("int $55");` in `main()`. Because XINU's IDT only configures entries for 0–47, interrupt #55 lacks a valid entry. The CPU detects an invalid interrupt vector and raises a General Protection Fault #13.

I also see a trap message and a panic, and never return to `main()` to print anything afterward. This happens because XINU calls `set_evec(i, ...)` in a loop only up to `i < 48` in `evec.c`. The IDT has no entry for #55, so hardware triggers a GPF.

## Testing-

```
kprintf("Before int $55\n");
asm("int $55");
kprintf("After int $55\n");
```

## Output-

```
agraw185,Agrawal,Shivli,Spring,2025
Before int $55
Xinu trap!
exception 13 (general protection violation) currrpid 3 (Main process)
error code 000001ba (442)
CS EFC0008 eip 10281F
eflags 10246
register dump:
eax 00000001 (1)
ecx 00000000 (0)
edx 00116A20 (1141280)
ebx 00000000 (0)
esp 0EFC8FCC (251432908)
ebp 0EFC8FCC (251432908)
esi 00000000 (0)
edi 00000000 (0)
```

```
panic: Trap processing complete...
```

Currently, `pidt->igd_type = IGDT_TRAPG` value = 15 decimal, i.e. 0xF, which is a trap gate. A trap gate does not automatically disable external interrupts on entry. Instead, clock code in `clkdisp.S` explicitly does `cli` on entry and `sti` before returning.

So to change clock interrupt (#32) from trap gate to interrupt gate I changed only the clock interrupt vector #32 to be an interrupt gate type=0xE, or 14 decimal, which automatically disabled interrupts on entry and re-enabled on `iret`.

```
pidt->igd_type = IGDT_TRAPG;
if (xnum == 32) {
    pidt->igd_type = 0xE;
}
```

On rerunning the Lab 1.4.1 clock experiment, the clock still incremented normally, and timing was correct. The code in `clkdisp.S` that does `cli/sti` became somewhat redundant because hardware was already disabling interrupts upon entry. Given the same results as before, confirm that leaving `cli/sti` calls in `clkdisp.S` does not conflict with the CPU's automatic disabling in an interrupt gate. The kernel code remains correct and functional.

Then, I saw that selectors for DS -> 0x10 and SS -> 0x18 in the GDT. This is different from Linux/Windows, which unify them in 32-bit mode.

So I replaced the line that loads SS with the same selector as DS in `start.S`:

```
asm
movl  $0x10, %eax
movw  %ax, %ds
movw  %ax, %es
movw  %ax, %fs
movw  %ax, %gs
//movl $0x18, %eax
movw  %ax, %ss
```

And left the separate stack descriptor in the GDT unused. Now on rerunning the same Lab 1.4.1 clock test, the system behaved exactly the same. Because XINU does not rely on hardware segmentation for separate user vs. kernel stacks, unifying SS and DS caused no visible differences.

### 4.2.1

For this part, I started with creating system/getppidtrap.c and defined it in prototypes.h 'pid32 getppidtrap(pid32 pid);'. Then I implemented getppidtrap() to function as the system call API. It accepts one argument, a pid32 value. Then loads the system call number (21) into register EDX. On executing the trap instruction int \$35, it switches to kernel mode. After the trap completes, it obtains the return value from EAX and returns it.

First implementation-

```
#include <xinu.h>
#include "prototypes.h"

pid32 getppidtrap(pid32 pid)
{
    pid32 retval; (Variable to store the return value)
    asm volatile (
        "movl $21, %%edx\n\t" (Place system call number 21 into EDX)
        "int $35\n\t" (software interrupt)
        "movl %%eax, %0\n\t" (return value from EAX into retval)
        : "=r" (retval)
        :
        : "edx", "eax", "ecx", "ebx", "esi", "edi" (clobbered registers)
    );
    return retval;
}
```

Then I realized that my clobber list conflicts with the chosen input/output registers, because it marked EAX as both clobbered and the register output. So I specified a more precise output constraint-

```
#include <xinu.h>
#include "prototypes.h"
pid32 getppidtrap(pid32 pid)
{
    pid32 retval;
    asm volatile (
        "movl $21, %%edx\n\t"
        "int $35\n\t"
        : "=a" (retval)
        :
        : "edx", "ecx", "ebx", "esi", "edi", "memory"
    );
    return retval;
}
```

This approach came after a bit of brainstorming, where the line "=a"(retval) tells GCC the output is in register EAX. So there's no separate instruction inside the assembly block like movl %%eax, %0.

Because now EAX is an output, I don't list eax as a clobbered register. Instead, let GCC handle it via the output operand. I still declare EDX as clobbered because the assembly modifies it. And also add memory to the clobber if the kernel code might change memory.

Testing for this-

```
process main(void)

{
    pid32 ppid;
    kprintf("Invoking trapped getppid system call ... \n");
    ppid = getppidtrap(getpid());
    kprintf("Trapped getppid returned: %d\n", ppid);
    return OK;
}
```

Initial output-

```
agraw185,Agrawal,Shivli,Spring,2025
Invoking trapped getppid system call ...
Xinu trap!
exception 35 currpri 3 (Main process)
CS EFC0008 eip 101BFD
eflags 206
register dump:
eax 00000003 (3)
ecx 00000000 (0)
edx 00000015 (21)
ebx 00000000 (0)
esp 0EFC8F9C (251432860)
ebp 0EFC8F9C (251432860)
esi 00000000 (0)
edi 00000000 (0)
```

```
panic: Trap processing complete...
```

This tells that the system call wraps getppidtrap() to do int \$35, but Xtrap → trap(), which panics by default. So the code never actually calls getpid(). So I call a Dispatcher Function from \_Xint35. It reads the system call number from EDX, calls the corresponding internal kernel function getppid(), returns the result in EAX, and then irets.

```
.globl _Xint35
.extern syscalldisp

_Xint35:
    pushal
    call syscalldisp
    popal
    iret
```

Syscalldisp.c-

```
#include <xinu.h>
extern pid32 getppid(pid32);

void syscalldisp(void)
{
    uint32 scnum;
```

```

asm("movl %%edx, %0" : "=r" (scnum));

switch(scnum) {
case 21:
{
    pid32 arg;
    asm("movl 44(%%esp), %0" : "=r" (arg));
    pid32 ret = getppid(arg);
    asm("movl %0, %%eax" : : "r" (ret));
    break;
}

default:
    asm("movl $-1, %eax");
    break;
}
}

```

Now on testing again with the above main, I see-

```

Invoking trapped getppid system call...
Trapped getppid returned: 3

```

Which means the trapped system call is working. `getppidtrap()` executed in `$35`, the system call dispatcher at `_Xint35` recognized the system call number 21, fetched the argument from the user stack, called `getppid()`, and put its return value in EAX. Control returned to `getppidtrap()`, which printed the parent PID 3. No panic or trap message means successfully bypassing the old `jmp Xtrap` path and using the new dispatcher logic.

#### 4.2.2

For this the system call dispatcher, the second software layer I had to handle disabling interrupts on entry. Then, checking the system call number stored in register EDX. Also, extracting arguments from the user stack, and calling the corresponding upper-half kernel function, like `getppid()`. Finally, return the function's result in EAX and ensure the stack is in the correct state before executing `iret` to return to user mode since this ensures that system calls from user mode are correctly forwarded to their kernel-mode implementations. So I modified `_Xint35` in `intr.S` to handle system calls, which previously jumped to `Xtrap` causing a panic, with a custom system call dispatcher. `cli` is used to prevent unwanted stack shifts.

```

        .globl _Xint35
        .extern getppid
_Xint35:
    cli
    pushal
    cmpl $21, %edx
    je    .do_getppid
    movl  $-1, %eax
    jmp    .returnFromSyscall

.do_getppid:
    movl  44(%%esp), %ecx
    pushl %ecx
    call  getppid
    addl  $4, %esp

```

```
.returnFromSyscall:
    popal
    iret
```

Initially, I assumed the first argument passed was located at 44(%esp), but I saw incorrect results (parent PID = 512, 518, 544, etc.). This was because disable() was adding extra function call overhead, shifting the expected stack layout. To debug and correct the offset, I printed out values at different stack locations-

```
movl 36(%esp), %eax
pushl %eax
pushl $msg36      # Debug print: "36(%esp) = %d\n"
call kprintf
addl $8, %esp

movl 40(%esp), %eax
pushl %eax
pushl $msg40      # Debug print: "40(%esp) = %d\n"
call kprintf
addl $8, %esp

movl 44(%esp), %eax
pushl %eax
pushl $msg44      # Debug print: "44(%esp) = %d\n"
call kprintf
addl $8, %esp
```

#### Output-

```
36(%esp)= 251396104
40(%esp)= 518
44(%esp)= 0
48(%esp)= 0
Returned parent pid: 544
```

After testing, I confirmed that 44(%esp) was the correct location for the pid argument when cli was used instead of call disable().

#### Then on final testing-

```
process main(void)
{
    kprintf("Calling getppidtrap() ... \n");
    pid32 ppid = getppidtrap(getpid());
    kprintf("Returned parent pid: %d\n", ppid);
    return OK;
}
```

#### Output-

```
Calling getppidtrap() ...
Returned parent pid: 3
```

This confirmed that the system call successfully traps into \_Xint35 using int \$35 and the dispatcher correctly extracts the user argument from the stack. The dispatcher invokes getppid() and retrieves the correct result. And, EAX holds the correct parent PID (3) upon return.



Now, I had to support 0, 1, or 2 argument system calls as the current only supported system calls with a single argument like getpid(pid32 pid). So I checked the system call number stored in EDX and correctly branch to:

- .do\_getpid for 0 arguments.
- .do\_getppid for 1 argument.
- .do\_chprio for 2 arguments.

I started by finding that the first argument is located at 44(%esp) after pushal. And the second argument is at 48(%esp). The arguments are pushed right-to-left before calling the function. So after the function call, I adjusted the stack using addl \$X, %esp to remove the pushed arguments.

### Final implementation-

```
_Xint35:
    cli
    pushal
    cmpl    $20, %edx
    je      .do_getpid
    cmpl    $21, %edx    # check 2
    je      .do_getppid
    cmpl    $22, %edx    # check 3
    je      .do_chprio
    movl    $-1, %eax
    jmp     .returnFromSyscall
# 0 arguments (getpidtrap)
.do_getpid:
    call    getpid
    jmp     .returnFromSyscall
# 1 argument (getppidtrap)
.do_getppid:
    movl    44(%esp), %ecx
    pushl   %ecx
    call    getppid
    addl    $4, %esp
    jmp     .returnFromSyscall
# 2 arguments (chprio)
.do_chprio:
    movl    44(%esp), %ecx
    movl    48(%esp), %ebx
    pushl   %ebx
    pushl   %ecx
    call    chprio
    addl    $8, %esp
    jmp     .returnFromSyscall
.returnFromSyscall:
    popal
    iret
```

### Then I ran a final test-

```
process main(void)
{
    kprintf("Testing getpidtrap() ... \n");
    pid32 mypid = getpidtrap();
    kprintf("Returned pid: %d\n", mypid);

    kprintf("Testing getppidtrap() ... \n");
    pid32 parent = getppidtrap(getpid());
    kprintf("Returned parent pid: %d\n", parent);
```

```

    kprintf("Testing chpriotrap()...\n");
    pri16 old_prio = chpriotrap(getpid(), 5);
    kprintf("Returned old priority: %d\n", old_prio);

    return OK;
}

```

## Output-

```

agraw185,Agrawal,Shivli,Spring,2025
Testing getpidtrap() ...
Returned pid: 1
Testing getppidtrap() ...
Returned parent pid: 3
Testing chprio()...
Returned old priority: 5

```

This confirmed that the system calls correctly transition from user space to kernel space and back.

### 4.2.3

As stated and implemented above, I reused XINU's existing system calls—`getpid()`, `getppid()`, and `chprio()`, as upper-half kernel functions. These functions execute the service on behalf of user-mode applications that invoke system calls via `int $35`. The system call dispatcher (`_Xint35`) calls these upper-half functions while ensuring that `EAX` is not overwritten before executing `iret` and the CDECL calling convention is followed to correctly pass and return arguments.

As implemented above, I updated `_Xint35` to invoke `getpid()`, `getppid()`, and `chprio()` based on the system call number in `EDX` and each function executed in C and returned a value in `EAX`, which is preserved for the user wrapper. `_Xint35` also ensured that after a system call like `getppid()`, `EAX` is not altered before returning to user mode. The final `_Xint35` implementation correctly calls kernel functions and returns via `iret`, preserving results. I modified `getppidtrap()` to return from the trap, it copies the value in `EAX` to a local variable, which is then returned to the main.

In the `getppidtrap()` implementation, `EAX` is preserved by `_Xint35` and read back into `retval` by `getppidtrap()` and `retval` is returned to `main()`, completing the required flow.

```

#include <xinu.h>
#include "prototypes.h"
pid32 getppidtrap(pid32 pid)
{
    pid32 retval;
    asm volatile (
        "movl $21, %%edx      \n\t"
        "int $35              \n\t"
        : "=a" (retval)
        :
        : "edx", "ecx", "ebx", "esi", "edi", "memory"
    );
    return retval;
}

```

## 4.3

Trapped system call support for these XINU calls:

getpid() – 0 arguments – system call number 20

getppid() – 1 argument – system call number 21

chprio() – 2 arguments – system call number 22

Corresponding wrapper functions:

getpidtrap() in system/getpidtrap.c

getppidtrap() in system/getppidtrap.c

chpriotrap() in system/chpriotrap.c

In include/process.h I added-

```
#define SYSCALL_GETPID 20
#define SYSCALL_GETPPID 21
#define SYSCALL_CHPRIO 22
```

so \_Xint35 knows how to dispatch calls based on EDX.

## 4.4

As I have shown testing in all the above sections, I will conclude by highlighting that I tested the trapped system calls by calling-

```
pid32 mypid = getpidtrap();
pid32 parent = getppidtrap(mypid);
pri16 oldprio = chpriotrap(mypid, 5);
```

and verifying that mypid is a sensible PID & parent is the parent's PID & oldprio matches the process's old priority.

I confirmed that no panic or trap message appears, meaning \_Xint35 is correctly extracting arguments and returning EAX unmodified.

```
/* main.c - main */
#include <xinu.h>
process main(void)
{
    kprintf("Testing getpidtrap() ... \n");
    pid32 mypid = getpidtrap();
    kprintf("Returned pid: %d\n", mypid);
    kprintf("Testing getppidtrap() ... \n");
    pid32 parent = getppidtrap(getpid());
    kprintf("Returned parent pid: %d\n", parent);
    kprintf("Testing chprio() ... \n");
    pri16 old_prio = chpriotrap(getpid(), 5);
    kprintf("Returned old priority: %d\n", old_prio);
    return OK;
}
```

Output-

agrawal185,Agrawal,Shivli, Spring,2025

```
Testing getpidtrap() ...
Returned pid: 1
Testing getppidtrap() ...
Returned parent pid: 3
Testing chprio() ...
Returned old priority: 5
```

## Bonus-

For this, I extended the trapped system call implementation to support `suspend()`. `suspend()` is a system call that puts the current process into a suspended state, removing it from the ready queue and forcing a context switch via `resched()`. A suspended process remains inactive until it is explicitly resumed.

Then I implemented `suspendtrap()`, and assigned it system call number 23. On being called, the wrapper places the value 23 in register EDX and issues an `int $35` instruction. The system call dispatcher then recognizes this call and invokes the `suspend()` function as the upper-half kernel function. This function changes the process's state, removes it from the ready list, and returns the process's old priority code..

```
syscall suspendtrap(pid32 pid)
{
    syscall retval;
    asm volatile (
        "movl $23, %%edx      \n\t"
        "int $35              \n\t"
        : "=a" (retval)
        : "b" (pid)
        : "edx", "ecx", "esi", "edi", "memory"
    );

    return retval;
}
```

I extended the dispatcher in `_Xint35` to handle system call number 23. For a `suspend` call, the dispatcher extracts the single argument from the user stack and calls the `suspend()` function like `getppidtrap`. I made no changes to the `suspend` function. Instead, repurposed it as the upper-half kernel function for the trapped `suspend` system call. The dispatcher calls `suspend` with the target PID, and its return value is passed through to the `suspendtrap`.

This ensures that the dispatcher does not disturb the contents of EAX after the `suspend` returns so that `suspendtrap` can correctly retrieve the old priority value.

## To test-

```
process test_suspend(void)
{
    while (1) {
        kprintf("Running process %d\n", getpid());
        sleep(2);
    }
    return OK;
}

process main(void)
{
```

```
pid32 testpid = create(test_suspend, 8192, 20, "test_suspend", 0);
resume(testpid);
sleep(3);
status s = suspendtrap(testpid);
kprintf("Suspended process %d ⇒ status %d\n", testpid, s);
return OK;
}
```

I then concluded that if a process suspends itself or is the only runnable process, the scheduler might immediately reselect it. So the bonus problem implements that.